Similar model-call cost, very different success

The useful trade-off is higher success with fewer model calls.

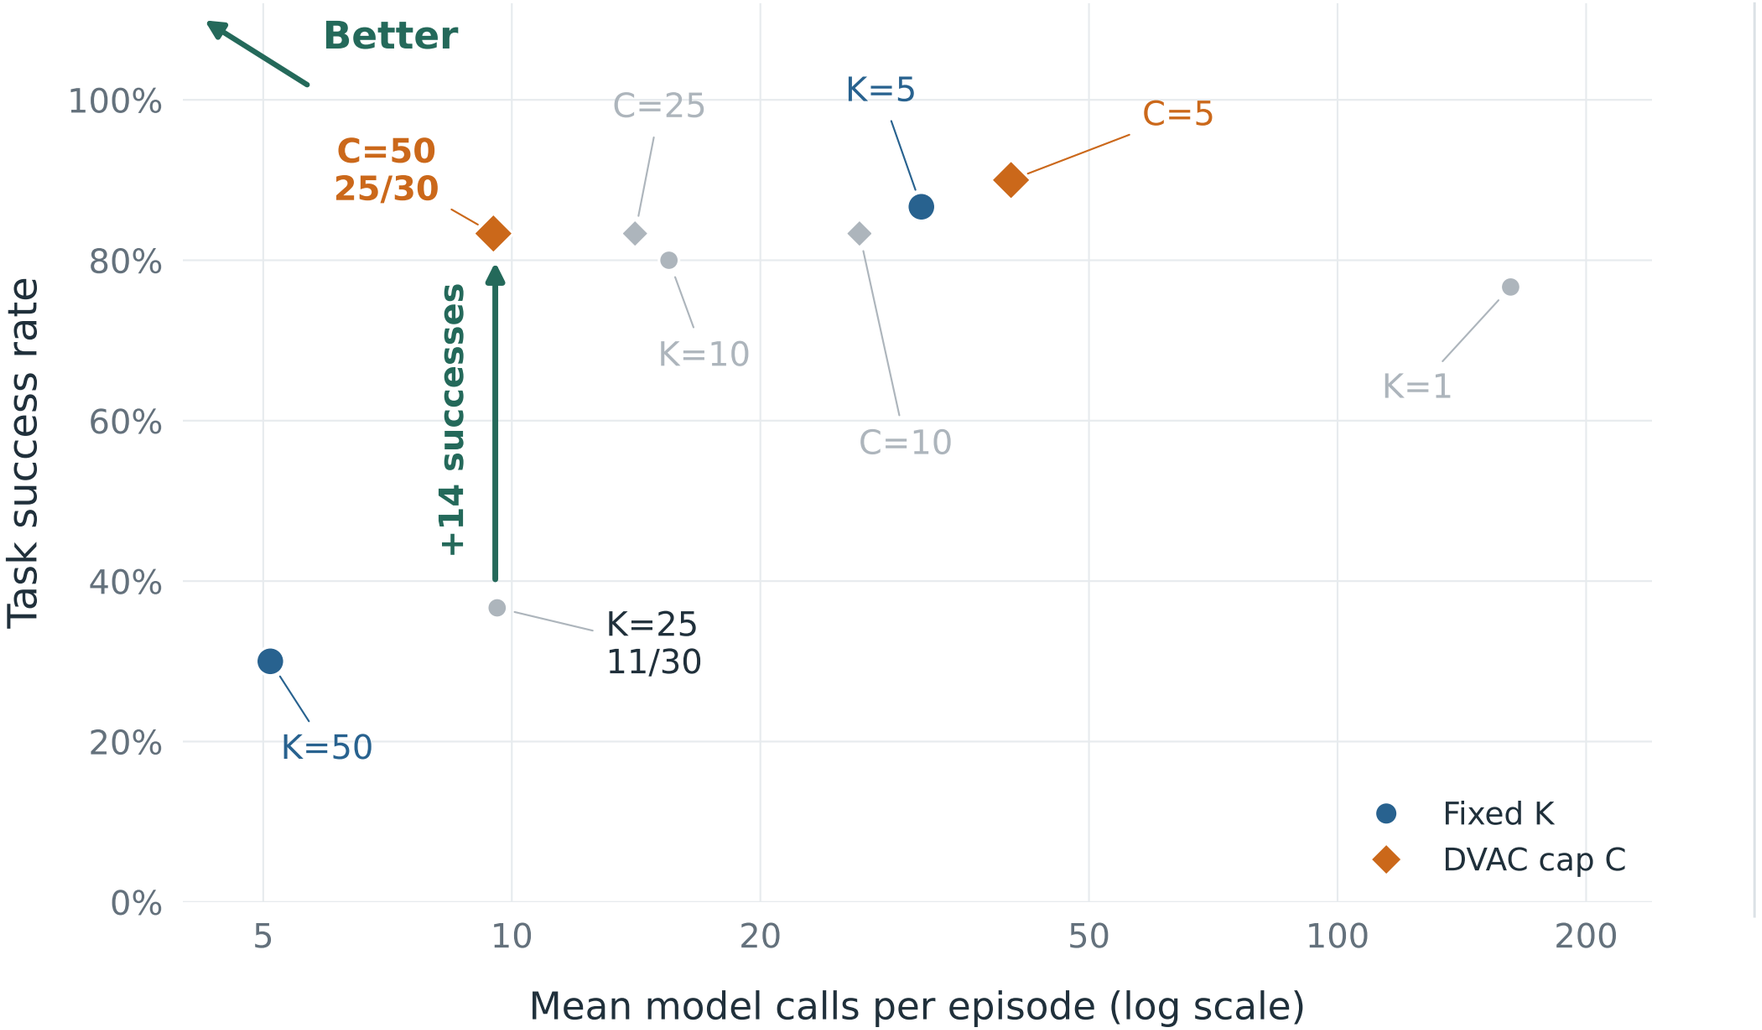

Almost the same cost

DVAC C=50

25/30 successes · 9.5 calls

Fixed K=25

11/30 successes · 9.6 calls

Higher observed success

Setting	Success	Calls
C=50	25/30	9.5
K=5	26/30	31.3
C=5	27/30	40.2

Colored points: best observed trade-offs. For each gray point, a lower-cost option did at least as well.

30 matched episodes per method. Calls include failures. These are sample results, without uncertainty intervals.

K is the fixed execution interval. C is the maximum adaptive interval. No single setting wins on both objectives.